

DIJKSTRA SHORTEST PATH ALGORITHM PROOF

Suppose we want to find the shortest path from a starting node s to all other nodes in a directed graph with no negative weights. Set S contains all the nodes that we have already been processed (meaning we have calculated their minimum distance from s), and the set Q contains the rest. Each time we process a node, we check all its neighbors to see whether we can update their distance (relaxing edge) or not.

Algorithm 1 Dijkstra

Require: Graph $G(V, E, w)$, $w : e \in E(G) \rightarrow R^+$

Ensure: $d[u]$: minimum distance from vertex s to all other vertices

```

1:  $d[s] \leftarrow 0$ 
2:  $d[v] \leftarrow \infty, \forall v \in V(G)$ 
3:  $U \leftarrow \emptyset$ 
4: while  $|U| < |V|$  do
5:    $u \leftarrow \operatorname{argmin}\{d[v], v \in V(G) \setminus U\}$ 
6:    $U \leftarrow U + \{u\}$ 
7:   for all  $v \in N(u)$  do
8:      $d[v] \leftarrow \min(d[v], d[u] + w(u, v))$ 
9:   end for
10: end while
11: return  $d$ 
```

Theorem 1. *At the end of the algorithm 1, $d[v]$ contains the minimum distance from s to every other $v \in V(G)$.*

We will prove it by induction. Let $d(s, v)$ be the real distance of s to any other vertex v . We have to show that $d[v] = d(s, v)$ for each vertex $v \in V(G)$. It's obvious that $d(s, s) = d[s] = 0$. Consider now any iteration where $d(s, v) = d[v]$ holds for all the previous vertices added to U and algorithm chooses v^* to add to U . We will show that $d(s, v^*) = d[v^*]$. Now let any path P from s to v^* with length equal to $d(s, v^*)$. Assume, for sake of contradiction, that $d[v^*] > d(s, v^*)$ (1). Because $s \in U$ and $v^* \notin U$ there exists at least one edge (x, y) on P that $x \in U$ and $y \notin U$ (could be $x = s$ or $y = v^*$). We pick the first such edge. By induction it follows that $d[x] = d(s, x)$. Providing that weights are non negative and P contains a $s - x$ path (plus the edge (x, y)), it holds that $d[x] + w(x, y) \leq d(s, v^*)$ (2). Moreover, we calculated $d[y]$, from line 8, as $d[y] = \min(d[y], d[x] + w(x, y))$ after putting x to U . This means that $d[y] \leq d[x] + w(x, y)$ (3). Hence, $d[y] \leq d(s, v^*)$ from (2) and (3). But this would mean that $d[y] < d[v^*]$. This is a contradiction, because we assumed that the algorithm picked v^* (line 5) at the beginning of the iteration and not y . See the next page for a possible part of an input graph. Complexity is $O(|V|^2)$ because we do $|V|$ iterations (line 4) and argmin takes $O(|V|)$ (line 5). We can do better ($O(|V| + |E| \log |E|)$) using a heap.

Green vertices are in U
Red vertices are not in U




Figure 1: An example of a path P . Note that it may leave U and enter U again.